







Imputation wall times from audited benchmark snapshots
n = 100,000 per dataset; m = 5; median with IQR across 5 repeats; 2 warmups excluded
Same time scale within each dataset; different scales across datasets. Only recorded methods are shown.



R reference bars are repeated from the native audit for context; reference and torch are different implementations.
R full pipeline includes preprocessing, bootstrap, draws, postprocessing and bridge costs; native Python omits the R bridge.
Suites ran in different wall-clock batches. Startup/import and post-run scoring are excluded. IQR is not a confidence interval.
Common host is caller-attested, not verified by this plotter. Optional ratios only pair same dtype inside the same route/suite.